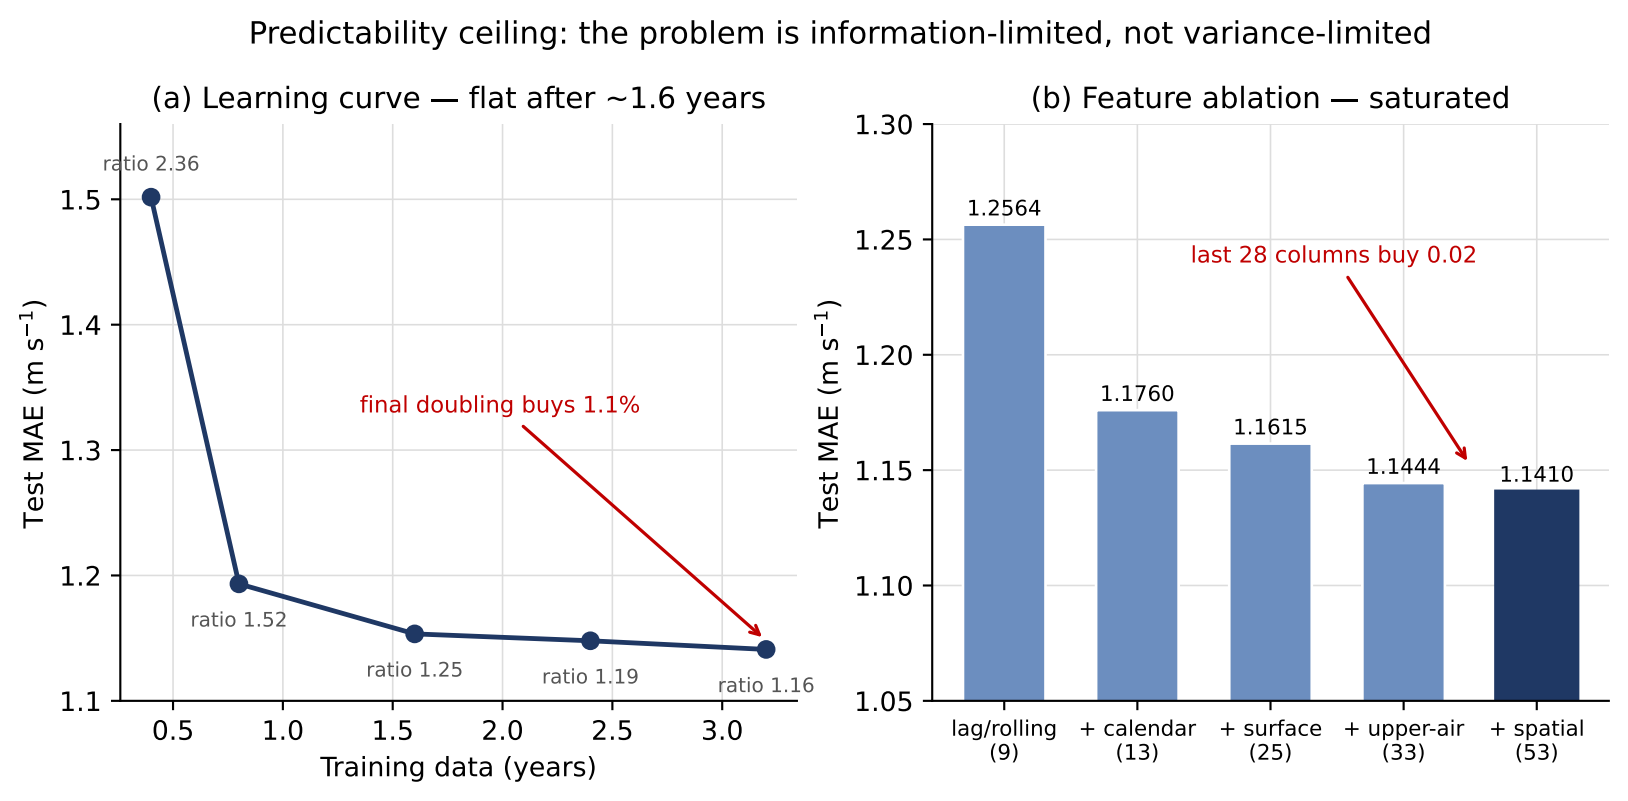
Predictability ceiling: the problem is information-limited, not variance-limited
(a) Learning curve — flat after ~1.6 years
Test MAE (m s^{-1})
Training data (years)
ratio 2.36
ratio 1.52
ratio 1.25
ratio 1.19
ratio 1.16
final doubling buys 1.1%
(b) Feature ablation — saturated
Test MAE (m s^{-1})
1.2564
1.1760
1.1615
1.1444
1.1410
last 28 columns buy 0.02
lag/rolling (9)
+ calendar (13)
+ surface (25)
+ upper-air (33)
+ spatial (53)