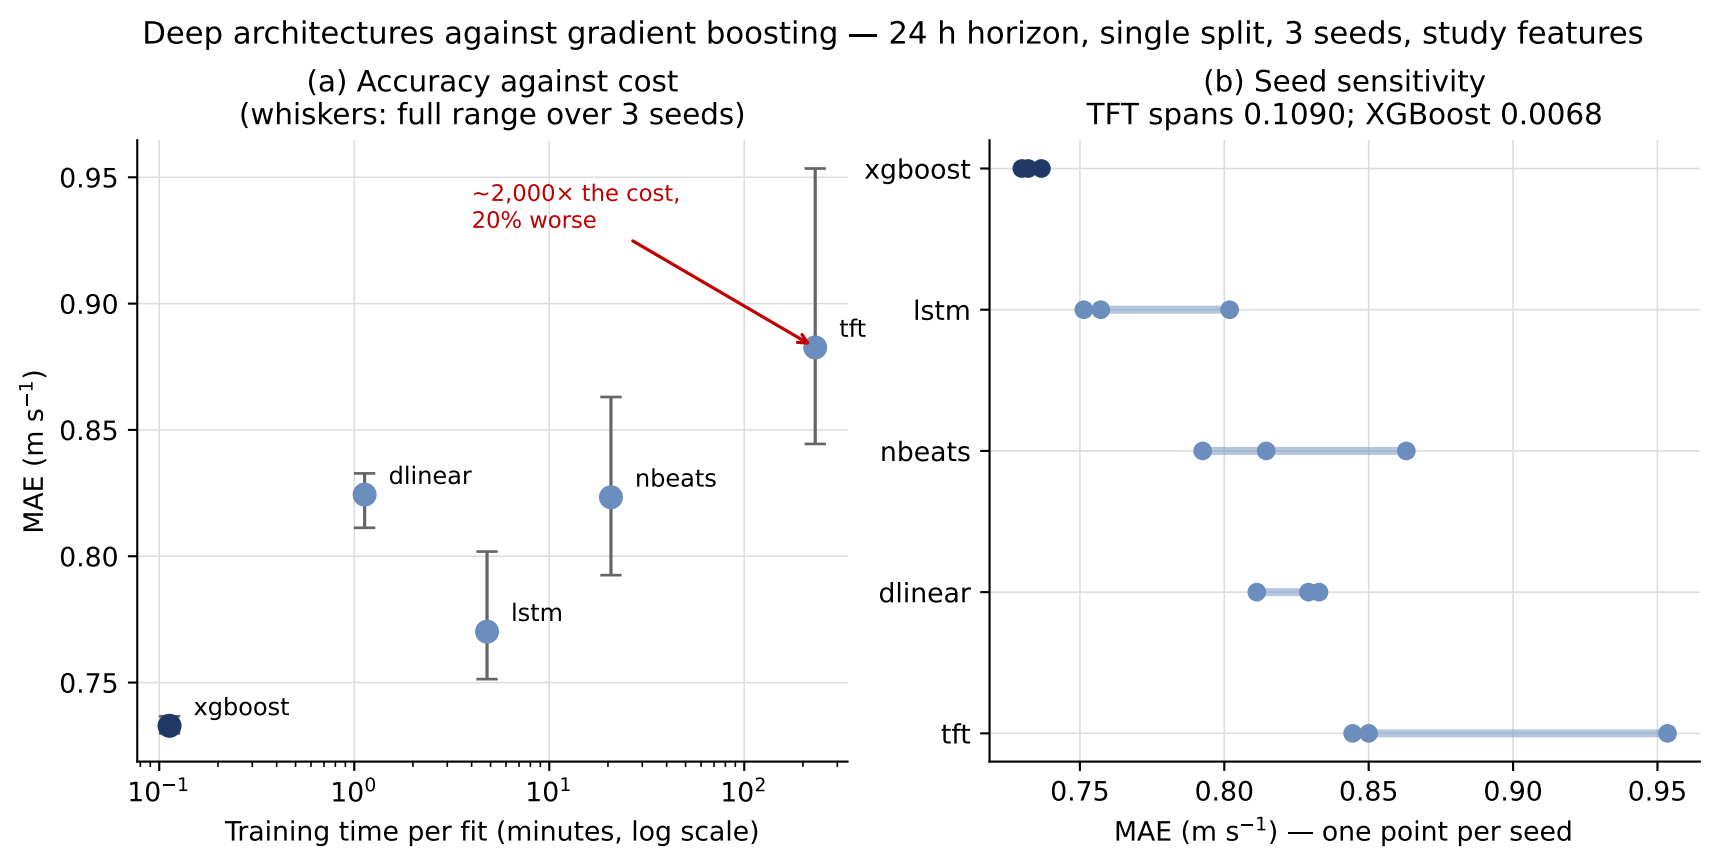
Deep architectures against gradient boosting — 24 h horizon, single split, 3 seeds, study features
(a) Accuracy against cost
(whiskers: full range over 3 seeds)
~2,000× the cost,
20% worse
xgboost
dlinear
lstm
nbeats
tft
MAE (m s^{-1})
Training time per fit (minutes, log scale)
(b) Seed sensitivity
TFT spans 0.1090; XGBoost 0.0068
MAE (m s^{-1}) — one point per seed